

[Configure SAP HANA 1.0 System Replication](#)[Create SAP HANA cluster resources](#)[Test the cluster setup](#)[Next steps](#)

For on-premises development, you can use either HANA System Replication or use shared storage to establish high availability for SAP HANA. On Azure virtual machines (VMs), HANA System Replication on Azure is currently the only supported high availability function. SAP HANA Replication consists of one primary node and at least one secondary node. Changes to the data on the primary node are replicated to the secondary node synchronously or asynchronously.


This article describes how to deploy and configure the virtual machines, install the cluster framework, and install and configure SAP HANA System Replication. In the example configurations, installation commands, instance number **03**, and HANA System ID **HN1** are used.


Read the following SAP Notes and papers first:

- SAP Note [1928533](#), which has:
 - The list of Azure VM sizes that are supported for the deployment of SAP software.
 - Important capacity information for Azure VM sizes.
 - The supported SAP software, and operating system (OS) and database combinations.
 - The required SAP kernel version for Windows and Linux on Microsoft Azure.
- SAP Note [2015553](#) lists the prerequisites for SAP-supported SAP software deployments in Azure.
- SAP Note [2205917](#) has recommended OS settings for SUSE Linux Enterprise Server for SAP Applications.
- SAP Note [1944799](#) has SAP HANA Guidelines for SUSE Linux Enterprise Server for SAP Applications.
- SAP Note [2178632](#) has detailed information about all of the monitoring metrics that are reported for SAP in Azure.
- SAP Note [2191498](#) has the required SAP Host Agent version for Linux in Azure.
- SAP Note [2243692](#) has information about SAP licensing on Linux in Azure.
- SAP Note [1984787](#) has general information about SUSE Linux Enterprise Server 12.
- SAP Note [1999351](#) has additional troubleshooting information for the Azure Enhanced Monitoring Extension for SAP.
- SAP Note [401162](#) has information on how to avoid "address already in use" when setting up HANA System Replication.
- [SAP Community WIKI](#) has all of the required SAP Notes for Linux.
- [SAP HANA Certified IaaS Platforms](#)
- [Azure Virtual Machines planning and implementation for SAP on Linux](#) guide.
- [Azure Virtual Machines deployment for SAP on Linux](#) (this article).
- [Azure Virtual Machines DBMS deployment for SAP on Linux](#) guide.
- [SUSE Linux Enterprise Server for SAP Applications 12 SP3 best practices guides](#)
 - Setting up an SAP HANA SR Performance Optimized Infrastructure (SLES for SAP Applications 12 SP1). The guide contains all of the required information to set up SAP HANA System Replication for on-premises development. Use this guide as a baseline.
 - Setting up an SAP HANA SR Cost Optimized Infrastructure (SLES for SAP Applications 12 SP1)

Overview

To achieve high availability, SAP HANA is installed on two virtual machines. The data is replicated by using HANA System Replication.





SAP HANA System Replication setup uses a dedicated virtual hostname and virtual IP addresses. On Azure, a load balancer is required to use a virtual IP address. The following list shows the configuration of the load balancer:

- Front-end configuration: IP address 10.0.0.13 for hn1-db
- Back-end configuration: Connected to primary network interfaces of all virtual machines that should be part of HANA System Replication
- Probe Port: Port 62503
- Load-balancing rules: 30313 TCP, 30315 TCP, 30317 TCP

Deploy for Linux

The resource agent for SAP HANA is included in SUSE Linux Enterprise Server for SAP Applications. The Azure Marketplace contains an image for SUSE Linux Enterprise Server for SAP Applications 12 that you can use to deploy new virtual machines.

Deploy with a template

You can use one of the quickstart templates that are on GitHub to deploy all the required resources. The template deploys the virtual machines, the load balancer, the availability set, and so on. To deploy the template, follow these steps:

1. Open the [database template](#) or the [converged template](#) on the Azure portal. The database template creates the load-balancing rules for a database only. The converged template also creates the load-balancing rules for an ASCS/SCS and ERS (Linux only) instance. If you plan to install an SAP NetWeaver-based system and you want to install the ASCS/SCS instance on the same machines, use the [converged template](#).
2. Enter the following parameters:
 - **Sap System ID:** Enter the SAP system ID of the SAP system you want to install. The ID is used as a prefix for the resources that are deployed.
 - **Stack Type:** (This parameter is applicable only if you use the converged template.) Select the SAP NetWeaver stack type.
 - **Os Type:** Select one of the Linux distributions. For this example, select **SLES 12**.
 - **Db Type:** Select **HANA**.
 - **Sap System Size:** Enter the number of SAPS that the new system is going to provide. If you're not sure how many SAPS the system requires, ask your SAP Technology Partner or System Integrator.
 - **System Availability:** Select **HA**.
 - **Admin Username and Admin Password:** A new user is created that can be used to sign in to the machine.
 - **New Or Existing Subnet:** Determines whether a new virtual network and subnet should be created or an existing subnet used. If you already have a virtual network that's connected to your on-premises network, select **Existing**.
 - **Subnet ID:** If you want to deploy the VM into an existing VNet where you have a subnet defined the VM should be assigned to, name the ID of that specific subnet. The ID usually looks like `/subscriptions/<subscription ID>/resourceGroups/<resource group name>/providers/Microsoft.Network/virtualNetworks/<virtual network name>/subnets/<subnet name>`.

Manual deployment

❗ Important

Make sure that the OS you select is SAP certified for SAP HANA on the specific VM types you are using. The list of SAP HANA certified VM types and OS releases for those can be looked up in [SAP HANA Certified IaaS Platforms](#). Make sure to click into the details of the VM type

listed to get the complete list of SAP HANA supported OS releases for the specific VM type

1. Create a resource group.
2. Create a virtual network.
3. Create an availability set.
 - Set the max update domain.
4. Create a load balancer (internal).
 - Select the virtual network created in step 2.
5. Create virtual machine 1.
 - Use a SLES4SAP image in the Azure gallery that is supported for SAP HANA on the VM type you selected.
 - Select the availability set created in step 3.
6. Create virtual machine 2.
 - Use a SLES4SAP image in the Azure gallery that is supported for SAP HANA on the VM type you selected.
 - Select the availability set created in step 3.
7. Add data disks.
8. Configure the load balancer. First, create a front-end IP pool:
 - a. Open the load balancer, select **frontend IP pool**, and select **Add**.
 - b. Enter the name of the new front-end IP pool (for example, **hana-frontend**).
 - c. Set the **Assignment to Static** and enter the IP address (for example, **10.0.0.13**).
 - d. Select **OK**.
 - e. After the new front-end IP pool is created, note the pool IP address.
9. Next, create a back-end pool:
 - a. Open the load balancer, select **backend pools**, and select **Add**.
 - b. Enter the name of the new back-end pool (for example, **hana-backend**).
 - c. Select **Add a virtual machine**.
 - d. Select the availability set created in step 3.
 - e. Select the virtual machines of the SAP HANA cluster.
 - f. Select **OK**.
10. Next, create a health probe:
 - a. Open the load balancer, select **health probes**, and select **Add**.
 - b. Enter the name of the new health probe (for example, **hana-hp**).
 - c. Select **TCP** as the protocol and port 62503. Keep the **Interval** value set to 5, and the **Unhealthy threshold** value set to 2.
 - d. Select **OK**.
11. For SAP HANA 1.0, create the load-balancing rules:
 - a. Open the load balancer, select **load balancing rules**, and select **Add**.
 - b. Enter the name of the new load balancer rule (for example, hana-lb-30315).
 - c. Select the front-end IP address, the back-end pool, and the health probe that you created earlier (for example, **hana-frontend**).
 - d. Keep the **Protocol** set to **TCP**, and enter port 30315.
 - e. Increase the **idle timeout** to 30 minutes.
 - f. Make sure to **enable Floating IP**.
 - g. Select **OK**.
 - h. Repeat these steps for port 30317.
12. For SAP HANA 2.0, create the load-balancing rules for the system database:
 - a. Open the load balancer, select **load balancing rules**, and select **Add**.
 - b. Enter the name of the new load balancer rule (for example, hana-lb-30313).
 - c. Select the front-end IP address, the back-end pool, and the health probe that you created earlier (for example, **hana-frontend**).
 - d. Keep the **Protocol** set to **TCP**, and enter port 30313.
 - e. Increase the **idle timeout** to 30 minutes.
 - f. Make sure to **enable Floating IP**.

- f. Make sure to **enable floating IP**.
- g. Select **OK**.
- h. Repeat these steps for port **30314**.

13. For SAP HANA 2.0, first create the load-balancing rules for the tenant database:

- a. Open the load balancer, select **load balancing rules**, and select **Add**.
- b. Enter the name of the new load balancer rule (for example, hana-lb-**30340**).
- c. Select the frontend IP address, backend pool, and health probe you created earlier (for example, **hana-frontend**).
- d. Keep the **Protocol** set to **TCP**, and enter port **30340**.
- e. Increase the **idle timeout** to 30 minutes.
- f. Make sure to **enable Floating IP**.
- g. Select **OK**.
- h. Repeat these steps for ports **30341** and **30342**.

For more information about the required ports for SAP HANA, read the chapter [Connections to Tenant Databases](#) in the [SAP HANA Tenant Databases](#) guide or [SAP Note 2388694](#).

❗ Important

Do not enable TCP timestamps on Azure VMs placed behind Azure Load Balancer. Enabling TCP timestamps will cause the health probes to fail. Set parameter **net.ipv4.tcp_timestamps** to **0**. For details see [Load Balancer health probes](#). See also SAP note [2382421](#).

Create a Pacemaker cluster

Follow the steps in [Setting up Pacemaker on SUSE Linux Enterprise Server in Azure](#) to create a basic Pacemaker cluster for this HANA server. You can use the same Pacemaker cluster for SAP HANA and SAP NetWeaver (A)SCS.

Install SAP HANA

The steps in this section use the following prefixes:

- **[A]**: The step applies to all nodes.
- **[1]**: The step applies to node 1 only.
- **[2]**: The step applies to node 2 of the Pacemaker cluster only.

1. **[A]** Set up the disk layout: **Logical Volume Manager (LVM)**.

We recommend that you use LVM for volumes that store data and log files. The following example assumes that the virtual machines have four data disks attached that are used to create two volumes.

List all of the available disks:

	Copy
<pre>ls /dev/disk/azure/scsi1/lun*</pre>	

Example output:

	Copy
<pre>/dev/disk/azure/scsi1/lun0 /dev/disk/azure/scsi1/lun1 /dev/disk/azure/scsi1/lun2 /dev/disk/azure/scsi1/lun3</pre>	

Create physical volumes for all of the disks that you want to use:

	Copy
<pre>sudo pvcreate /dev/disk/azure/scsi1/lun0 sudo pvcreate /dev/disk/azure/scsi1/lun1 sudo pvcreate /dev/disk/azure/scsi1/lun2 sudo pvcreate /dev/disk/azure/scsi1/lun3</pre>	

Create a volume group for the data files. Use one volume group for the log files and one for the shared directory of SAP HANA:

```
sudo vgcreate vg_hana_data_HN1 /dev/disk/azure/scsi1/lun0 /dev/disk/azure/scsi1/lun1
sudo vgcreate vg_hana_log_HN1 /dev/disk/azure/scsi1/lun2
sudo vgcreate vg_hana_shared_HN1 /dev/disk/azure/scsi1/lun3
```

Create the logical volumes. A linear volume is created when you use `lvcreate` without the `-i` switch. We suggest that you create a striped volume for better I/O performance, where the `-i` argument should be the number of the underlying physical volume. In this document, two physical volumes are used for the data volume, so the `-i` switch argument is set to **2**. One physical volume is used for the log volume, so no `-i` switch is explicitly used. Use the `-i` switch and set it to the number of the underlying physical volume when you use more than one physical volume for each data, log, or shared volumes.

```
sudo lvcreate -i 2 -l 100%FREE -n hana_data vg_hana_data_HN1
sudo lvcreate -l 100%FREE -n hana_log vg_hana_log_HN1
sudo lvcreate -l 100%FREE -n hana_shared vg_hana_shared_HN1
sudo mkfs.xfs /dev/vg_hana_data_HN1/hana_data
sudo mkfs.xfs /dev/vg_hana_log_HN1/hana_log
sudo mkfs.xfs /dev/vg_hana_shared_HN1/hana_shared
```

Create the mount directories and copy the UUID of all of the logical volumes:

```
sudo mkdir -p /hana/data/HN1
sudo mkdir -p /hana/log/HN1
sudo mkdir -p /hana/shared/HN1
# Write down the ID of /dev/vg_hana_data_HN1/hana_data, /dev/vg_hana_log_HN1/hana_log, and
/dev/vg_hana_shared_HN1/hana_shared
sudo blkid
```

Create `fstab` entries for the three logical volumes:

```
sudo vi /etc/fstab
```

Insert the following line in the `/etc/fstab` file:

```
/dev/disk/by-uuid/<UUID of /dev/mapper/vg_hana_data_HN1-hana_data> /hana/data/HN1 xfs defaults,nofail 0 2
/dev/disk/by-uuid/<UUID of /dev/mapper/vg_hana_log_HN1-hana_log> /hana/log/HN1 xfs defaults,nofail 0 2
/dev/disk/by-uuid/<UUID of /dev/mapper/vg_hana_shared_HN1-hana_shared> /hana/shared/HN1 xfs defaults,nofail 0 2
```

Mount the new volumes:

```
sudo mount -a
```

2. [A] Set up the disk layout: **Plain Disks**.

For demo systems, you can place your HANA data and log files on one disk. Create a partition on `/dev/disk/azure/scsi1/lun0` and format it with xfs:

```
sudo sh -c 'echo -e "\n\n\n\n\n\n\n" | fdisk /dev/disk/azure/scsi1/lun0'
sudo mkfs.xfs /dev/disk/azure/scsi1/lun0-part1

# Write down the ID of /dev/disk/azure/scsi1/lun0-part1
sudo /sbin/blkid
```

```
sudo vi /etc/fstab
```

Insert this line in the /etc/fstab file:

```
/dev/disk/by-uuid/<UUID> /hana xfs defaults,nofail 0 2
```

Create the target directory and mount the disk:

```
sudo mkdir /hana
sudo mount -a
```

3. **[A]** Set up host name resolution for all hosts.

You can either use a DNS server or modify the /etc/hosts file on all nodes. This example shows you how to use the /etc/hosts file. Replace the IP address and the hostname in the following commands:

```
sudo vi /etc/hosts
```

Insert the following lines in the /etc/hosts file. Change the IP address and hostname to match your environment:

```
10.0.0.5 hn1-db-0
10.0.0.6 hn1-db-1
```

4. **[A]** Install the SAP HANA high availability packages:

```
sudo zypper install SAPHanaSR
```

To install SAP HANA System Replication, follow chapter 4 of the [SAP HANA SR Performance Optimized Scenario guide](#).

1. **[A]** Run the **hdbicm** program from the HANA DVD. Enter the following values at the prompt:

- Choose installation: Enter **1**.
- Select additional components for installation: Enter **1**.
- Enter Installation Path [/hana/shared]: Select Enter.
- Enter Local Host Name [...]: Select Enter.
- Do you want to add additional hosts to the system? (y/n) [n]: Select Enter.
- Enter SAP HANA System ID: Enter the SID of HANA, for example: **HN1**.
- Enter Instance Number [00]: Enter the HANA Instance number. Enter **03** if you used the Azure template or followed the manual deployment section of this article.
- Select Database Mode / Enter Index [1]: Select Enter.
- Select System Usage / Enter Index [4]: Select the system usage value.
- Enter Location of Data Volumes [/hana/data/HN1]: Select Enter.
- Enter Location of Log Volumes [/hana/log/HN1]: Select Enter.
- Restrict maximum memory allocation? [n]: Select Enter.
- Enter Certificate Host Name For Host '...' [...]: Select Enter.
- Enter SAP Host Agent User (sapadm) Password: Enter the host agent user password.
- Confirm SAP Host Agent User (sapadm) Password: Enter the host agent user password again to confirm.
- Enter System Administrator (hdbadm) Password: Enter the system administrator password.
- Confirm System Administrator (hdbadm) Password: Enter the system administrator password again to confirm.
- Enter System Administrator Home Directory [/usr/sap/HN1/home]: Select Enter.
- Enter System Administrator Login Shell [/bin/sh]: Select Enter.

- Enter System Administrator User ID [1001]: Select Enter.
- Enter ID of User Group (sapsys) [79]: Select Enter.
- Enter Database User (SYSTEM) Password: Enter the database user password.
- Confirm Database User (SYSTEM) Password: Enter the database user password again to confirm.
- Restart system after machine reboot? [n]: Select Enter.
- Do you want to continue? (y/n): Validate the summary. Enter **y** to continue.

2. [A] Upgrade the SAP Host Agent.

Download the latest SAP Host Agent archive from the [SAP Software Center](#) and run the following command to upgrade the agent. Replace the path to the archive to point to the file that you downloaded:

```
sudo /usr/sap/hostctrl/exe/saphostexec -upgrade -archive <path to SAP Host Agent SAR>
```

 Copy

Configure SAP HANA 2.0 System Replication

The steps in this section use the following prefixes:

- **[A]**: The step applies to all nodes.
- **[1]**: The step applies to node 1 only.
- **[2]**: The step applies to node 2 of the Pacemaker cluster only.

1. [1] Create the tenant database.

If you're using SAP HANA 2.0 or MDC, create a tenant database for your SAP NetWeaver system. Replace **NW1** with the SID of your SAP system.

Execute the following command as <hanasid>adm :

```
hdbsql -u SYSTEM -p "passwd" -i 03 -d SYSTEMDB 'CREATE DATABASE NW1 SYSTEM USER PASSWORD "passwd"'
```

 Copy

2. [1] Configure System Replication on the first node:

Back up the databases as <hanasid>adm:

```
hdbsql -d SYSTEMDB -u SYSTEM -p "passwd" -i 03 "BACKUP DATA USING FILE ('initialbackupSYS')"
```

```
hdbsql -d HN1 -u SYSTEM -p "passwd" -i 03 "BACKUP DATA USING FILE ('initialbackupHN1')"
```

```
hdbsql -d NW1 -u SYSTEM -p "passwd" -i 03 "BACKUP DATA USING FILE ('initialbackupNW1')"
```

 Copy

Copy the system PKI files to the secondary site:

```
scp /usr/sap/HN1/SYS/global/security/rsecssfs/data/SSFS_HN1.DAT hn1-db-1:/usr/sap/HN1/SYS/global/security/rsecssfs/data/
```

```
scp /usr/sap/HN1/SYS/global/security/rsecssfs/key/SSFS_HN1.KEY hn1-db-1:/usr/sap/HN1/SYS/global/security/rsecssfs/key/
```

 Copy

Create the primary site:

```
hdbnsutil -sr_enable --name=SITE1
```

 Copy

3. [2] Configure System Replication on the second node:

Register the second node to start the system replication. Run the following command as <hanasid>adm :

 Copy

```
sapcontrol -nr 03 -function StopWait 600 10
hdbnsutil -sr_register --remoteHost=hn1-db-0 --remoteInstance=03 --replicationMode=sync --name=SITE2
```

Configure SAP HANA 1.0 System Replication

The steps in this section use the following prefixes:

- **[A]**: The step applies to all nodes.
- **[1]**: The step applies to node 1 only.
- **[2]**: The step applies to node 2 of the Pacemaker cluster only.

1. **[1]** Create the required users.

Run the following command as root. Make sure to replace bold strings (HANA System ID **HN1** and instance number **03**) with the values of your SAP HANA installation:

```
PATH="$PATH:/usr/sap/HN1/HDB03/exe"
hdbsql -u system -i 03 'CREATE USER hdbhasync PASSWORD "passwd"'
hdbsql -u system -i 03 'GRANT DATA ADMIN TO hdbhasync'
hdbsql -u system -i 03 'ALTER USER hdbhasync DISABLE PASSWORD LIFETIME'
```

 Copy

2. **[A]** Create the keystore entry.

Run the following command as root to create a new keystore entry:

```
PATH="$PATH:/usr/sap/HN1/HDB03/exe"
hdbuserstore SET hdbhaloc localhost:30315 hdbhasync passwd
```

 Copy

3. **[1]** Back up the database.

Back up the databases as root:

```
PATH="$PATH:/usr/sap/HN1/HDB03/exe"
hdbsql -d SYSTEMDB -u system -i 03 "BACKUP DATA USING FILE ('initialbackup')"
```

 Copy

If you use a multi-tenant installation, also back up the tenant database:

```
hdbsql -d HN1 -u system -i 03 "BACKUP DATA USING FILE ('initialbackup')"
```

 Copy

4. **[1]** Configure System Replication on the first node.

Create the primary site as <hanasid>adm :

```
su - hdbadm
hdbnsutil -sr_enable --name=SITE1
```

 Copy

5. **[2]** Configure System Replication on the secondary node.

Register the secondary site as <hanasid>adm:

```
sapcontrol -nr 03 -function StopWait 600 10
hdbnsutil -sr_register --remoteHost=hn1-db-0 --remoteInstance=03 --replicationMode=sync --name=SITE2
```

 Copy

Create SAP HANA cluster resources

First, create the HANA topology. Run the following commands on one of the Pacemaker cluster nodes:

Copy

```
sudo crm configure property maintenance-mode=true

# Replace the bold string with your instance number and HANA system ID

sudo crm configure primitive rsc_SAPHanaTopology_HN1_HDB03 ocf:suse:SAPHanaTopology \
  operations \ $id="rsc_sap2_HN1_HDB03-operations" \
  op monitor interval="10" timeout="600" \
  op start interval="0" timeout="600" \
  op stop interval="0" timeout="300" \
  params SID="HN1" InstanceNumber="03"

sudo crm configure clone cln_SAPHanaTopology_HN1_HDB03 rsc_SAPHanaTopology_HN1_HDB03 \
  meta is-managed="true" clone-node-max="1" target-role="Started" interleave="true"
```

Next, create the HANA resources:

Copy

```
# Replace the bold string with your instance number, HANA system ID, and the front-end IP address of the Azure load balancer.

sudo crm configure primitive rsc_SAPHana_HN1_HDB03 ocf:suse:SAPHana \
  operations \ $id="rsc_sap_HN1_HDB03-operations" \
  op start interval="0" timeout="3600" \
  op stop interval="0" timeout="3600" \
  op promote interval="0" timeout="3600" \
  op monitor interval="60" role="Master" timeout="700" \
  op monitor interval="61" role="Slave" timeout="700" \
  params SID="HN1" InstanceNumber="03" PREFER_SITE_TAKEOVER="true" \
  DUPLICATE_PRIMARY_TIMEOUT="7200" AUTOMATED_REGISTER="false"

sudo crm configure ms msl_SAPHana_HN1_HDB03 rsc_SAPHana_HN1_HDB03 \
  meta is-managed="true" notify="true" clone-max="2" clone-node-max="1" \
  target-role="Started" interleave="true"

sudo crm configure primitive rsc_ip_HN1_HDB03 ocf:heartbeat:IPaddr2 \
  meta target-role="Started" is-managed="true" \
  operations \ $id="rsc_ip_HN1_HDB03-operations" \
  op monitor interval="10s" timeout="20s" \
  params ip="10.0.0.13"

sudo crm configure primitive rsc_nc_HN1_HDB03 anything \
  params binfile="/usr/bin/nc" cmdline_options="-l -k 62503" \
  op monitor timeout=20s interval=10 depth=0

sudo crm configure group g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 rsc_nc_HN1_HDB03

sudo crm configure colocation col_saphana_ip_HN1_HDB03 4000: g_ip_HN1_HDB03:Started \
  msl_SAPHana_HN1_HDB03:Master

sudo crm configure order ord_SAPHana_HN1_HDB03 Optional: cln_SAPHanaTopology_HN1_HDB03 \
  msl_SAPHana_HN1_HDB03

# Clean up the HANA resources. The HANA resources might have failed because of a known issue.
sudo crm resource cleanup rsc_SAPHana_HN1_HDB03

sudo crm configure property maintenance-mode=false
sudo crm configure rsc_defaults resource-stickiness=1000
sudo crm configure rsc_defaults migration-threshold=5000
```

Make sure that the cluster status is ok and that all of the resources are started. It's not important on which node the resources are running.

Copy

```
sudo crm_mon -r
```

```
# Online: [ hn1-db-0 hn1-db-1 ]
#
# Full list of resources:
#
# stonith-sbd      (stonith:external/sbd): Started hn1-db-0
# rsc_st_azure     (stonith:fence_azure_arm):      Started hn1-db-1
# Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
#   Started: [ hn1-db-0 hn1-db-1 ]
# Master/Slave Set: msl_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
#   Masters: [ hn1-db-0 ]
#   Slaves: [ hn1-db-1 ]
# Resource Group: g_ip_HN1_HDB03
#   rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2):      Started hn1-db-0
#   rsc_nc_HN1_HDB03 (ocf::heartbeat:anything):      Started hn1-db-0
```

Test the cluster setup

This section describes how you can test your setup. Every test assumes that you are root and the SAP HANA master is running on the **hn1-db-0** virtual machine.

Test the migration

Before you start the test, make sure that Pacemaker does not have any failed action (via `crm_mon -r`), there are no unexpected location constraints (for example leftovers of a migration test) and that HANA is sync state, for example with `SAPHanaSR-showAttr`:

Copy

```
hn1-db-0:~ # SAPHanaSR-showAttr

Global cib-time
-----
global Mon Aug 13 11:26:04 2018

Hosts      clone_state lpa_hn1_lpt node_state op_mode  remoteHost  roles                                score site  srmode
sync_state version                vhost
-----
hn1-db-0 PROMOTED    1534159564 online    logreplay nws-hana-vm-1 4:P:master1:master:worker:master 150  SITE1 sync  PRIM
2.00.030.00.1522209842 nws-hana-vm-0
hn1-db-1 DEMOTED     30      online    logreplay nws-hana-vm-0 4:S:master1:master:worker:master 100  SITE2 sync  SOK
2.00.030.00.1522209842 nws-hana-vm-1
```

You can migrate the SAP HANA master node by executing the following command:

Copy

```
crm resource migrate msl_SAPHana_HN1_HDB03 hn1-db-1
```

If you set `AUTOMATED_REGISTER="false"`, this sequence of commands should migrate the SAP HANA master node and the group that contains the virtual IP address to hn1-db-1.

Once the migration is done, the `crm_mon -r` output looks like this

Copy

```
Online: [ hn1-db-0 hn1-db-1 ]

Full list of resources:

stonith-sbd      (stonith:external/sbd): Started hn1-db-1
Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: msl_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-1 ]
  Stopped: [ hn1-db-0 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2):      Started hn1-db-1
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything):      Started hn1-db-1
```

Failed Actions:

```
* rsc_SAPHana_HN1_HDB03_start_0 on hn1-db-0 'not running' (7): call=84, status=complete, exitreason='none',
  last-rc-change='Mon Aug 13 11:31:37 2018', queued=0ms, exec=2095ms
```

The SAP HANA resource on hn1-db-0 fails to start as secondary. In this case, configure the HANA instance as secondary by executing this command:



```
su - hn1adm

# Stop the HANA instance just in case it is running
hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> sapcontrol -nr 03 -function StopWait 600 10
hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 --replicationMode=sync --
name=SITE1
```

The migration creates location constraints that need to be deleted again:



```
# Switch back to root and clean up the failed state
exit
hn1-db-0:~ # crm resource unmigrate msl_SAPHana_HN1_HDB03
```

You also need to clean up the state of the secondary node resource:



```
hn1-db-0:~ # crm resource cleanup msl_SAPHana_HN1_HDB03 hn1-db-0
```

Monitor the state of the HANA resource using `crm_mon -r`. Once HANA is started on hn1-db-0, the output should look like this



```
Online: [ hn1-db-0 hn1-db-1 ]

Full list of resources:

stonith-sbd      (stonith:external/sbd): Started hn1-db-1
Clone Set: cIn_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: msl_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-1 ]
  Slaves: [ hn1-db-0 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1
```

Test the Azure fencing agent (not SBD)

You can test the setup of the Azure fencing agent by disabling the network interface on the hn1-db-0 node:



```
sudo ifdown eth0
```

The virtual machine should now restart or stop depending on your cluster configuration. If you set the `stonith-action` setting to off, the virtual machine is stopped and the resources are migrated to the running virtual machine.

After you start the virtual machine again, the SAP HANA resource fails to start as secondary if you set `AUTOMATED_REGISTER="false"`. In this case, configure the HANA instance as secondary by executing this command:



```
su - hn1adm
```

```
# Stop the HANA instance just in case it is running
sapcontrol -nr 03 -function StopWait 600 10
hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 --replicationMode=sync --name=SITE1

# Switch back to root and clean up the failed state
exit
crm resource cleanup msl_SAPHana_HN1_HDB03 hn1-db-0
```

Test SBD fencing

You can test the setup of SBD by killing the inquisitor process.



```
hn1-db-0:~ # ps aux | grep sbd
root      1912  0.0  0.0  85420 11740 ?        SL   12:25   0:00 sbd: inquisitor
root      1929  0.0  0.0  85456 11776 ?        SL   12:25   0:00 sbd: watcher: /dev/disk/by-id/scsi-
360014056f268462316e4681b704a9f73 - slot: 0 - uuid: 7b862dba-e7f7-4800-92ed-f76a4e3978c8
root      1930  0.0  0.0  85456 11776 ?        SL   12:25   0:00 sbd: watcher: /dev/disk/by-id/scsi-
360014059bc9ea4e4bac4b18808299aaf - slot: 0 - uuid: 5813ee04-b75c-482e-805e-3b1e22ba16cd
root      1931  0.0  0.0  85456 11776 ?        SL   12:25   0:00 sbd: watcher: /dev/disk/by-id/scsi-
36001405b8ddd44eb3647908def6621c - slot: 0 - uuid: 986ed8f8-947d-4396-8aec-b933b75e904c
root      1932  0.0  0.0  90524 16656 ?        SL   12:25   0:00 sbd: watcher: Pacemaker
root      1933  0.0  0.0 102708 28260 ?        SL   12:25   0:00 sbd: watcher: Cluster
root      13877 0.0  0.0   9292  1572 pts/0    S+   12:27   0:00 grep sbd

hn1-db-0:~ # kill -9 1912
```

Cluster node hn1-db-0 should be rebooted. The Pacemaker service might not get started afterwards. Make sure to start it again.

Test a manual failover

You can test a manual failover by stopping the `pacemaker` service on the hn1-db-0 node:



```
service pacemaker stop
```

After the failover, you can start the service again. If you set `AUTOMATED_REGISTER="false"`, the SAP HANA resource on the hn1-db-0 node fails to start as secondary. In this case, configure the HANA instance as secondary by executing this command:



```
service pacemaker start
su - hn1adm

# Stop the HANA instance just in case it is running
sapcontrol -nr 03 -function StopWait 600 10
hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 --replicationMode=sync --name=SITE1

# Switch back to root and clean up the failed state
exit
crm resource cleanup msl_SAPHana_HN1_HDB03 hn1-db-0
```

SUSE tests

Important

Make sure that the OS you select is SAP certified for SAP HANA on the specific VM types you are using. The list of SAP HANA certified VM types and OS releases for those can be looked up in [SAP HANA Certified IaaS Platforms](#). Make sure to click into the details of the VM type listed to get the complete list of SAP HANA supported OS releases for the specific VM type


Run all test cases that are listed in the SAP HANA SR Performance Optimized Scenario or SAP HANA SR Cost Optimized Scenario guide, depending on your use case. You can find the guides on the [SLES for SAP best practices page](#).

The following tests are a copy of the test descriptions of the SAP HANA SR Performance Optimized Scenario SUSE Linux Enterprise Server for SAP Applications 12 SP1 guide. For an up-to-date version, always also read the guide itself. Always make sure that HANA is in sync before starting the test and also make sure that the Pacemaker configuration is correct.


In the following test descriptions we assume `PREFER_SITE_TAKEOVER="true"` and `AUTOMATED_REGISTER="false"`. NOTE: The following tests are designed to be run in sequence and depend on the exit state of the preceding tests.

1. TEST 1: STOP PRIMARY DATABASE ON NODE 1

Resource state before starting the test:


	
<pre>Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-0] Slaves: [hn1-db-1] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0</pre>	

Run the following commands as `<hanasid>adm` on node `hn1-db-0`:


	
<pre>hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> HDB stop</pre>	

Pacemaker should detect the stopped HANA instance and failover to the other node. Once the failover is done, the HANA instance on node `hn1-db-0` is stopped because Pacemaker does not automatically register the node as HANA secondary.

Run the following commands to register node `hn1-db-0` as secondary and cleanup the failed resource.


	
<pre>hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 -- replicationMode=sync --name=SITE1 # run as root hn1-db-0:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-0</pre>	

Resource state after the test:

	
<pre>Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-1] Slaves: [hn1-db-0] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1</pre>	

2. TEST 2: STOP PRIMARY DATABASE ON NODE 2

Resource state before starting the test:

	
<pre>Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-1] Slaves: [hn1-db-0] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1</pre>	

Run the following commands as <hanasid>adm on node hn1-db-1:

Copy

```
hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> HDB stop
```

Pacemaker should detect the stopped HANA instance and failover to the other node. Once the failover is done, the HANA instance on node hn1-db-1 is stopped because Pacemaker does not automatically register the node as HANA secondary.

Run the following commands to register node hn1-db-1 as secondary and cleanup the failed resource.

Copy

```
hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-0 --remoteInstance=03 --
replicationMode=sync --name=SITE2

# run as root
hn1-db-1:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-1
```

Resource state after the test:

Copy

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

3. TEST 3: CRASH PRIMARY DATABASE ON NODE

Resource state before starting the test:

Copy

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

Run the following commands as <hanasid>adm on node hn1-db-0:

Copy

```
hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> HDB kill-9
```

Pacemaker should detect the killed HANA instance and failover to the other node. Once the failover is done, the HANA instance on node hn1-db-0 is stopped because Pacemaker does not automatically register the node as HANA secondary.


Run the following commands to register node hn1-db-0 as secondary and cleanup the failed resource.

Copy

```
hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 --
replicationMode=sync --name=SITE1


# run as root
hn1-db-0:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-0
```

Resource state after the test:


		
Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-1] Slaves: [hn1-db-0] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1		

4. TEST 4: CRASH PRIMARY DATABASE ON NODE 2

Resource state before starting the test:


		
Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-1] Slaves: [hn1-db-0] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1		

Run the following commands as <hanasid>adm on node hn1-db-1:


		
<pre>hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> HDB kill-9</pre>		

Pacemaker should detect the killed HANA instance and failover to the other node. Once the failover is done, the HANA instance on node hn1-db-1 is stopped because Pacemaker does not automatically register the node as HANA secondary.

Run the following commands to register node hn1-db-1 as secondary and cleanup the failed resource.


		
<pre>hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-0 --remoteInstance=03 -- replicationMode=sync --name=SITE2 # run as root hn1-db-1:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-1</pre>		

Resource state after the test:

		
Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03] Started: [hn1-db-0 hn1-db-1] Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03] Masters: [hn1-db-0] Slaves: [hn1-db-1] Resource Group: g_ip_HN1_HDB03 rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0 rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0		

5. TEST 5: CRASH PRIMARY SITE NODE (NODE 1)

Resource state before starting the test:

		
Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]		

```

Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
Masters: [ hn1-db-0 ]
Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0

```

Run the following commands as root on node hn1-db-0:

```
hn1-db-0:~ # echo 'b' > /proc/sysrq-trigger
```

Pacemaker should detect the killed cluster node and fence the node. Once the node is fenced, Pacemaker will trigger a takeover of the HANA instance. When the fenced node is rebooted, Pacemaker will not start automatically.

Run the following commands to start Pacemaker, clean the SBD messages for node hn1-db-0, register node hn1-db-0 as secondary, and cleanup the failed resource.

```

# run as root
# list the SBD device(s)
hn1-db-0:~ # cat /etc/sysconfig/sbd | grep SBD_DEVICE=
# SBD_DEVICE="/dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116;/dev/disk/by-id/scsi-36001405034a84428af24ddd8c3a3e9e1;/dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3"

hn1-db-0:~ # sbd -d /dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116 -d /dev/disk/by-id/scsi-36001405034a84428af24ddd8c3a3e9e1 -d /dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3 message hn1-db-0 clear

hn1-db-0:~ # systemctl start pacemaker

# run as <hanasid>adm
hn1adm@hn1-db-0:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-1 --remoteInstance=03 --replicationMode=sync --name=SITE1

# run as root
hn1-db-0:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-0

```

Resource state after the test:

```

Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
Masters: [ hn1-db-1 ]
Slaves: [ hn1-db-0 ]
Resource Group: g_ip_HN1_HDB03
rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1
rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1

```

6. TEST 6: CRASH SECONDARY SITE NODE (NODE 2)

Resource state before starting the test:

```

Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
Masters: [ hn1-db-1 ]
Slaves: [ hn1-db-0 ]
Resource Group: g_ip_HN1_HDB03
rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-1
rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-1

```

Run the following commands as root on node hn1-db-1:



```
hn1-db-1:~ # echo 'b' > /proc/sysrq-trigger
```

Pacemaker should detect the killed cluster node and fence the node. Once the node is fenced, Pacemaker will trigger a takeover of the HANA instance. When the fenced node is rebooted, Pacemaker will not start automatically.

Run the following commands to start Pacemaker, clean the SBD messages for node hn1-db-1, register node hn1-db-1 as secondary, and cleanup the failed resource.



```
# run as root
# list the SBD device(s)
hn1-db-1:~ # cat /etc/sysconfig/sbd | grep SBD_DEVICE=
# SBD_DEVICE="/dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116;/dev/disk/by-id/scsi-36001405034a84428af24dd8c3a3e9e1;/dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3"

hn1-db-1:~ # sbd -d /dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116 -d /dev/disk/by-id/scsi-36001405034a84428af24dd8c3a3e9e1 -d /dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3 message hn1-db-1 clear

hn1-db-1:~ # systemctl start pacemaker

# run as <hanasid>adm
hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> hdbnsutil -sr_register --remoteHost=hn1-db-0 --remoteInstance=03 --replicationMode=sync --name=SITE2

# run as root
hn1-db-1:~ # crm resource cleanup msl_SAPHana_HN1_HDB03 hn1-db-1
```

Resource state after the test:



```
Clone Set: cIn_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: msl_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

7. TEST 7: STOP THE SECONDARY DATABASE ON NODE 2

Resource state before starting the test:



```
Clone Set: cIn_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: msl_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

Run the following commands as <hanasid>adm on node hn1-db-1:



```
hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> HDB stop
```

Pacemaker will detect the stopped HANA instance and mark the resource as failed on node hn1-db-1. Pacemaker should automatically restart the HANA instance. Run the following command to clean up the failed state.



```
# run as root
hn1-db-1:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-1
```

Resource state after the test:

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

8. TEST 8: CRASH THE SECONDARY DATABASE ON NODE 2

Resource state before starting the test:

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

Run the following commands as <hanasid>adm on node hn1-db-1:

```
hn1adm@hn1-db-1:/usr/sap/HN1/HDB03> HDB kill-9
```

Pacemaker will detect the killed HANA instance and mark the resource as failed on node hn1-db-1. Run the following command to clean up the failed state. Pacemaker should then automatically restart the HANA instance.

```
# run as root
hn1-db-1:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-1
```

Resource state after the test:

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
  Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
  Masters: [ hn1-db-0 ]
  Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
  rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
  rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0
```

9. TEST 9: CRASH SECONDARY SITE NODE (NODE 2) RUNNING SECONDARY HANA DATABASE

Resource state before starting the test:

```
Clone Set: c1n_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
```

```

Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
Masters: [ hn1-db-0 ]
Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0

```

Run the following commands as root on node hn1-db-1:

```
hn1-db-1:~ # echo b > /proc/sysrq-trigger
```

Pacemaker should detect the killed cluster node and fence the node. When the fenced node is rebooted, Pacemaker will not start automatically.

Run the following commands to start Pacemaker, clean the SBD messages for node hn1-db-1, and cleanup the failed resource.

```

# run as root
# list the SBD device(s)
hn1-db-1:~ # cat /etc/sysconfig/sbd | grep SBD_DEVICE=
# SBD_DEVICE="/dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116;/dev/disk/by-id/scsi-36001405034a84428af24dd8c3a3e9e1;/dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3"

hn1-db-1:~ # sbd -d /dev/disk/by-id/scsi-36001405772fe8401e6240c985857e116 -d /dev/disk/by-id/scsi-36001405034a84428af24dd8c3a3e9e1 -d /dev/disk/by-id/scsi-36001405cdd5ac8d40e548449318510c3 message hn1-db-1 clear

hn1-db-1:~ # systemctl start pacemaker

hn1-db-1:~ # crm resource cleanup ms1_SAPHana_HN1_HDB03 hn1-db-1

```

Resource state after the test:

```

Clone Set: cln_SAPHanaTopology_HN1_HDB03 [rsc_SAPHanaTopology_HN1_HDB03]
Started: [ hn1-db-0 hn1-db-1 ]
Master/Slave Set: ms1_SAPHana_HN1_HDB03 [rsc_SAPHana_HN1_HDB03]
Masters: [ hn1-db-0 ]
Slaves: [ hn1-db-1 ]
Resource Group: g_ip_HN1_HDB03
rsc_ip_HN1_HDB03 (ocf::heartbeat:IPaddr2): Started hn1-db-0
rsc_nc_HN1_HDB03 (ocf::heartbeat:anything): Started hn1-db-0

```

Next steps

- [Azure Virtual Machines planning and implementation for SAP](#)
- [Azure Virtual Machines deployment for SAP](#)
- [Azure Virtual Machines DBMS deployment for SAP](#)
- To learn how to establish high availability and plan for disaster recovery of SAP HANA on Azure (large instances), see [SAP HANA \(large instances\) high availability and disaster recovery on Azure](#)

Is this page helpful?

Yes No

